Mischievous Sibling's Grid World

Aykut C. Satici[a], Gokhan Atinc[b]

[a]*Boise State University, Mechanical and Biomedical Engineering, Boise, 83725, Idaho, USA*
[b]*Mathworks Inc., 1 Lakeside Campus Drive, Natick, 01760, Massachusetts, USA*

Abstract

This is a technical note solving an interesting question posed by a friend and colleague, Gökhan Atınç. The question starts out exactly like the well-known Grid World problem, but it has a twist. We solve the problem in two different ways: first using a direct probability calculation and then using reinforcement learning on a judiciously designed Markov Decision Process. The results are compared and discussed.

Keywords: Probability, Expectations, Reinforcement learning, Gymnasium, Policy iteration

1. Introduction

The question is interesting and the solutions are just as interesting. Let us delve right into it.

2. Problem Statement

Alice has a 5×5 chess board as shown in Figure 1, where each square is described by its coordinates $(i, j) \in \{0, 1, 2, 3, 4\}^2$. Her remote-controlled robot starts on the square $(2, 2)$. The remote-controller has 4 buttons that can move the robot east, north, west, and south by 1 square with each press. Bob, her mischievous sibling, has tampered with the wiring of these buttons. Assuming that the buttons are fixed after having shuffled once, what is the expected value of the minimum number of times Alice needs to press the buttons in order to move the robot to the goal square $(4, 4)$?

Figure 1: Schematic of the problem.

For future reference, if we linearly index the squares of the chessboard from 0 to 24 in a row-major fashion, then the mapping from the linear index to the (i, j) coordinates and back is given by the following bijections:

$$5j + i = k \leftrightarrow \left(i = k - 5 \left\lfloor \frac{k}{5} \right\rfloor, j = \left\lfloor \frac{k}{5} \right\rfloor \right) \qquad (1)$$

3. Solutions

We present two solutions to the problem: a probabilistic solution and a reinforcement learning (RL) solution.

3.1. Probability Solution

We invoke a celebrated theorem from probability theory that will help us solve the problem. This theorem is called the *law of total expectation* or *tower rule of probability theory* in the literature Bertsekas and Tsitsiklis (2002). In this section, we use the correspondence $\{1, 2, 3, 4, 5\} \mapsto \{0, 1, 2, 3, 4\}$ and $\{A, B, C, D, E\} \mapsto \{0, 1, 2, 3, 4\}$ for legacy reasons.

Theorem 1 (Tower rule). *Let X be a random variable and $\{A_i\}_i^m$ be a finite partition of the sample space. Then,*

$$\mathbb{E}[X] = \sum_i^m \mathbb{E}[X \mid A_i] \, \mathbb{P}(A_i).$$

We denote the square on which the robot resides at the k^{th} step by s_k and the goal square by g. Consider the following family of random variables.

$$C_n := \sum_{k=n+1}^{\infty} r_k, \qquad r_k = \begin{cases} 1 & \text{if } s_k \neq g \\ 0 & \text{if } s_k = g \end{cases}. \qquad (2)$$

The sum in this definition is well-defined because once the robot reaches g, all the remaining r_k's take the value zero. We are being asked the value of $\mathbb{E}[C_0]$. To that end, we will use the following repercussion of definition (2).

$$C_m = \sum_{i=m+1}^{n} r_i + C_n, \qquad m \leq n.$$

Email addresses: aykutsatici@boisestate.edu (Aykut C. Satici), gokhan.atinc@mathworks.com (Gokhan Atinc)

Using the properties of expectation, if $s_i \neq g$ for $m < i < n$, we can deduce the following result.

$$\mathbb{E}[C_m] = \sum_{i=m+1}^{n} \mathbb{E}[r_i] + \mathbb{E}[C_n] = n - m + \mathbb{E}[C_n].$$

Using the properties of expectation, if $s_i \neq g$ for $m < i < n$, we can deduce the following result.

$$\mathbb{E}[C_m] = \sum_{i=m+1}^{n} \mathbb{E}[r_i] + \mathbb{E}[C_n] = n - m + \mathbb{E}[C_n].$$

Similar identities to the above expression hold for conditional expectations as well. Using the tower rule 1, we start to compute.

$$\mathbb{E}[C_0] = \mathbb{E}[C_0 \mid s_1 \in \{C4, D3\}] \underbrace{\mathbb{P}(s_1 \in \{C4, D3\})}_{=\frac{1}{2}} \\ + \mathbb{E}[C_0 \mid s_1 \in \{B3, C2\}] \underbrace{\mathbb{P}(s_1 \in \{B3, C2\})}_{=\frac{1}{2}}. \quad (3)$$

From now on, we will assume $s_1 = C4$ in the first term on the right-hand side and $s_1 = C2$ in the second. Notice that, because of the inherent symmetry of the problem, while the state $s_1 = D3$ is equivalent to $s_1 = C4$; $s_1 = B3$ is equivalent to $s_1 = C2$.

Computation of the first conditional expectation in (3). First of all, we consider the first expectation term on the right-hand side of equation (3). If at the first step, we found ourselves in square $C4$, that means we made progress towards the goal. We can keep this greedy motion for one more turn since it will get us closer to the goal. Hence, $s_2 = C5$, with a cost of $r_2 = 1$. Since at step 2, we hit the south wall, we need to change the button to press. We have 3 options of equal likelihood. Choosing one will take us either to $D5$, back to $C4$, or to $B5$. Hence,

$$\mathbb{E}[C_0 \mid s_1 = C4] = \overbrace{r_1 + r_2}^{=2} \\ + \underbrace{\mathbb{E}[C_2 \mid s_1 = C4, s_3 = D5]}_{=2} \underbrace{\mathbb{P}(s_3 = D5 \mid s_1 = C4)}_{=\frac{1}{3}} \\ + \mathbb{E}[C_2 \mid s_1 = s_3 = C4] \underbrace{\mathbb{P}(s_3 = C4 \mid s_1 = C4)}_{=\frac{1}{3}} \quad (4) \\ + \mathbb{E}[C_2 \mid s_1 = C4, s_3 = B5] \underbrace{\mathbb{P}(s_3 = B5 \mid s_1 = C4)}_{=\frac{1}{3}}.$$

We delve further into the computation of the two expected values in equation (4), whose values are not immediately apparent. If at step 3 we found ourselves back on $C4$, we got farther from the goal, but we can repeat our previous move to get back to $C5$ on step 4. This maneuver costs us $r_3 + r_4 = 2$ units. Now, we know what directions two of the keys map to. The remaining

expectations are computed below.

$$\mathbb{E}[C_2 \mid s_1 = s_3 = C4] = \underbrace{r_3 + r_4}_{=2} \\ + \underbrace{\mathbb{E}[C_4 \mid s_1 = s_3 = C4, s_5 = D5]}_{=2} \underbrace{\mathbb{P}(s_5 = D5 \mid s_1 = s_3 = C4)}_{=\frac{1}{2}} \\ + \underbrace{\mathbb{E}[C_4 \mid s_1 = s_3 = C4, s_5 = B5]}_{=4} \underbrace{\mathbb{P}(s_5 = B5 \mid s_1 = s_3 = C4)}_{=\frac{1}{2}}.$$

There is only one expected value that we have left to compute in equation (4). Again, if we found ourselves on the $B5$ square at step 3, we have moved farther away from the goal, but do now know how to get back to $C5$ in this case. Hence, we need to try out the final key to figure out the full mapping. The expectations are computed below.

$$\mathbb{E}[C_2 \mid s_1 = C4, s_3 = B5] = \\ \underbrace{\mathbb{E}[C_2 \mid s_1 = C4, s_3 = B5, s_4 = C5]}_{=4} \\ \times \underbrace{\mathbb{P}(s_4 = C5 \mid s_1 = C4, s_3 = B5)}_{=\frac{1}{2}} \\ + \underbrace{\mathbb{E}[C_2 \mid s_1 = C4, s_3 = B5, s_4 = B4]}_{=6} \\ \times \underbrace{\mathbb{P}(s_4 = B4 \mid s_1 = C4, s_3 = B5)}_{=\frac{1}{2}}.$$

The computations above allows us to determine the first expected value in equation (3) using equation (4): $\boxed{\mathbb{E}[C_0 \mid s_1 = C4] = 6}$.

Computation of the second conditional expectation in (3). We use similar techniques to compute the second expected value on the right-hand side of equation (3). If on step 2, we find ourselves on square $B2$, then, we have identified all actions that take us farther from the goal. The two remaining actions both get us closer to the goal, from which we are 8 steps away.

$$\mathbb{E}[C_0 \mid s_1 = C2] = \\ \underbrace{\mathbb{E}[C_0 \mid s_1 = C2, s_2 = B2]}_{=8} \underbrace{\mathbb{P}(s_2 = B2 \mid s_1 = C2)}_{=\frac{1}{3}} \\ + \mathbb{E}[C_0 \mid s_1 = C2, s_2 = C3] \underbrace{\mathbb{P}(s_2 = C3 \mid s_1 = C2)}_{=\frac{1}{3}} \quad (5) \\ + \mathbb{E}[C_0 \mid s_1 = C2, s_2 = D2] \underbrace{\mathbb{P}(s_2 = D2 \mid s_1 = C2)}_{=\frac{1}{3}}.$$

Once more, we utilize the tower rule to compute the expected values in equation (5) whose values are not immediately apparent. If at step 3, we end up on square $C3$, we can keep repeating this move until we reach $C5$ at step 4, incurring a cost of $r_1 + r_2 + r_3 + r_4 = 4$ units. From here, the only two possibilities is we get to square $D5$ or $B5$ on step 5. The remaining

expectations are computed below.

$$\mathbb{E}[C_0 \mid s_1 = C2, s_2 = C3] = \overbrace{r_1 + r_2 + r_3 + r_4}^{=4}$$
$$+ \underbrace{\mathbb{E}[C_4 \mid s_1 = C2, s_2 = C3, s_5 = D5]}_{=2}$$
$$\times \underbrace{\mathbb{P}(s_5 = D5 \mid s_1 = C2, s_2 = C3)}_{=\frac{1}{2}}$$
$$+ \underbrace{\mathbb{E}[C_4 \mid s_1 = C2, s_2 = C3, s_5 = B5]}_{=4}$$
$$\times \underbrace{\mathbb{P}(s_5 = B5 \mid s_1 = C2, s_2 = C3)}_{=\frac{1}{2}}.$$

There is only one expected value that we have left to compute in equation (5). If we find ourselves on square $D2$ at step 2, we can repeat this direction to reach $E2$ at step 3, incurring a cost of $r_1 + r_2 + r_3 = 3$ units. From here, we can either go to $E3$ or $D2$ with the remaining keys. The expectations are computed below.

$$\mathbb{E}[C_0 \mid s_1 = C2, s_2 = D2] = \overbrace{r_1 + r_2 + r_3}^{=3}$$
$$+ \underbrace{\mathbb{E}[C_3 \mid s_1 = C2, s_2 = D2, s_4 = E3]}_{=3}$$
$$\times \underbrace{\mathbb{P}(s_4 = E3 \mid s_1 = C2, s_2 = D2)}_{=\frac{1}{2}}$$
$$+ \underbrace{\mathbb{E}[C_3 \mid s_1 = C2, s_2 = s_4 = D2]}_{=5}$$
$$\times \underbrace{\mathbb{P}(s_4 = D2 \mid s_1 = C2, s_2 = D2)}_{=\frac{1}{2}}.$$

We can now use equation (5) in order to compute the second expected value on the right-hand side of equation (3): $\boxed{\mathbb{E}[C_0 \mid s_1 = C2] = \frac{22}{3}}$. We have computed every quantity that we are interested in. Now, we go back to equation (3):

$$\boxed{\begin{aligned}\mathbb{E}[C_0] &= \underbrace{\mathbb{E}\left[C_0 \mid s_1 \in \{C4, D3\}\right]}_{=6} \underbrace{\mathbb{P}(s_1 \in \{C4, D3\})}_{=\frac{1}{2}} \\ &+ \underbrace{\mathbb{E}\left[C_0 \mid s_1 \in \{B3, C2\}\right]}_{=\frac{22}{3}} \underbrace{\mathbb{P}(s_1 \in \{B3, C2\})}_{=\frac{1}{2}} \\ &= \frac{20}{3} = 6.\bar{6}.\end{aligned}}$$

3.2. RL solution

We want to be able to programmatically solve this toy problem so that we can find the solution for any given initial state. We can use reinforcement learning (RL) to solve this problem.

The reasoning goes as follows: we want to apply RL techniques, such as double Q-learning Morales (2020) to find the action-value function and from it the state-value function. However, in order to do this, we need to define the state, action, transition, and reward functions. The key to solving this problem

efficiently with RL methods is to observe that the state does not only comprise the position of the robot on the chessboard, but also the "world belief". Here, we define world belief as our current belief of mapping of the arrow keys to the actual directions.

Let $\underline{N} = \{0, 1, \ldots, N - 1\}$ denote the set of nonnegative integers ranging from 0 to $N - 1$. We model this problem as consisting of two RL subproblems: the first one is the sequential problem of finding the shortest route to the goal position, given the world belief, and the second is to find an appropriate world belief, which can be modeled as a 24-armed bandit problem.

We define the Markov decision processes (MDP) of the subproblems as follows:

3.2.1. State Spaces

The first subproblem has the state space $\mathcal{S}_1 = \underline{4} \times \underline{4}$, where the first factor stands for the column in which the robot is located, and the second stands for the row. The second subproblem is a bandit problem, with a singleton state, the world belief. The full problem may be considered to have the state space $\mathcal{S} = \mathcal{S}_1 \times \{1\} \cong \underline{4} \times \underline{4}$.

3.2.2. Action Spaces

The problem has the action space $\mathcal{A} = \underline{4} \times \underline{24}$. The first factor corresponds to the which key to be pressed given the world belief, which is an element of the second factor of the action space.

3.2.3. Transition function

The transition function is a deterministic function of the true world. The true world has a certain mapping between the keys and the directions, which is hidden from the agent. This key mapping is deterministic albeit hidden and the agent will move in the corresponding directions given the correct key mapping one hundred percent of the time.

3.2.4. Reward function

This function needs to be designed judiciously in order to give the agent the best chance to learn the optimal strategy. What worked well in our experiments is to assign a reward of -1 each time the agent moves into a nonterminal state (a state that is not G) for the grid world subproblem.

For the bandit subproblem, we define the reward as a sum of two terms. The first term is the difference between the value functions of the gridworld problem, given that the current world belief is correct. To explain, suppose that $w \in \underline{24}$ is the current world belief. We perform value iteration to find the value function of the grid world under this assumption and we assign the first term of the reward for the bandit as the value function evaluated at the grid world position at which we land minus the value function evaluated at the current grid world position.

It turns out that it is important to be able to distinguish between whether the desired movement direction is equal to the true direction of motion. For that reason, the second term that we add to the bandit reward signal -2 whenever the executed direction is different from the direction in which the agent expected the robot to move.

Figure 2: Running average number of steps taken to reach the goal over $10,000$ independent Monte Carlo simulation runs.

Our implementation Satici (2024) of this MDP is carried out using Gymnasium Towers et al. (2024), a Python library for modeling and RL.

3.2.5. Q-Learning

Once the MDP is defined, we apply Q-learning to find the optimal policy. For the grid world subproblem, this is the straightforward Q-learning algorithm that can be found in any RL textbook. For the bandit subproblem, the Q-function is learned by the iteration

$$Q_b[a] = Q_b[a] + (\texttt{reward} - Q_b[a])/\texttt{N}[a],$$

where $\texttt{N}[a]$ is the number of times action a has been taken for the bandit problem.

The way we set this problem up, allows us to use RL to learn to get to the goal position in one episode in the last number of moves. Of course, this is a stochastic process, depending on what the initial world belief is selected and what the true world is. We perform Monte Carlo simulations of this setup to determine the statistics, including the expectation of the minimum number of moves (which was solved analytically in Section 3.1).

4. Experiment

The simulation experiments we run in Gymnasium supports the theoretical computation performed in Section 3.1. The Monte Carlo simulations that we perform can yield even more results, such as what is the expected number of "bad" moves that the robot makes before it reaches the goal. Here, a bad move means that the agent moves away from the goal. We can further compute a histogram of the states visited, and the rewards collected on average.

The first result in Figure 2 shows the running average number of steps taken to reach the goal over $10,000$ independent Monte Carlo simulation runs. Figure 2 shows that the running average

converges to $6.\bar{6}$ just as the theoretical computation indicated in Section 3.1.

Figure 3 shows the histogram of number of steps it took to solve the environment over $10,000$ independent Monte Carlo simulation runs. We observe that it takes 4 steps to solve the environment $16.\bar{6}\%$ of the time, 6 steps $33.\bar{3}\%$ of the time, and 8 steps 50% of the time, yielding an expected value of

$$\frac{1}{6} \times 4 + \frac{1}{3} \times 6 + \frac{1}{2} \times 8 = \frac{20}{3} = 6.\bar{6}.$$


Figure 3: Histogram of number of steps taken to solve the environment over $10,000$ independent Monte Carlo simulation runs.

Figure 4 shows the frequency of visits of each state under the optimal policy deduced by Q-learning over $10,000$ independent Monte Carlo runs. Notice that the ratio of the goal state being visited to the number of runs is 1, meaning that all runs end at the goal state.


Figure 4: Ratio of times a given state is visited under the optimal policy over $10,000$ independent Monte Carlo simulation runs.

Figure 5 shows the scores each world belief achieves over simulation steps during a single run. At the end there are four

4

bandit arms(world beliefs) that have the same high score of 1, however, two of these bandits arms have been identified to swap the "down" and "right" actions. Since this swapping does not affect the performance of the agent, their score matches the other two top bandit arms which have not been differentiated.


Figure 5: Scores each world belief achieves over a single run sample.

5. Conclusions

Interesting problem solved.

References

Bertsekas, D.P., Tsitsiklis, J.N., 2002. Introduction to probability. volume 1. Athena Scientific Belmont, MA.

Morales, M., 2020. Grokking Deep Reinforcement Learning. Manning Publications.

Satici, A.C., 2024. Convoluted questions. URL: https://github.com/Symplectomorphism/miscellaneous/tree/master/dolambac.

Towers, M., Kwiatkowski, A., Terry, J., Balis, J.U., De Cola, G., Deleu, T., Goulão, M., Kallinteris, A., Krimmel, M., KG, A., et al., 2024. Gymnasium: A standard interface for reinforcement learning environments. arXiv preprint arXiv:2407.17032 .